

# COP 6726: Database Systems Implementation

## Spring 2018

### Weekly Assignment 5

#### Notes:

- Hashing Algorithm
- Let's start with the assumption that we have enough memory, if we don't we will deal with that later.
- Not having enough memory would mean we try external hashes because internal memory cannot hold all the information.
- How to we create these has functions?
- First consider how they are used, each of the tuple computes  $h(T)$  and then it is added to the corresponding bucket.
- Hash based schemes are little slower to build but are amazing are supporting lookups across the table.
- One of the things that became bad by switching to 64 bit architecture is that division started taking 80 cycles to complete compared to 40 on 32bit architecture.
- $h'(t) = h(t) \% B$
- Here B is the address size. Now modulus is what kills us with 80 cycles per instruction.
- So to avoid this massive waste, you try to find B which is in order of 2 and do bit shifting.
- Thus  $B=2^k$  and shift by k to divide it.
- Mercene primes  $2^k - 1$  are thus a really important factor and helps us choose B which is very good for our specific situations.
- Two different strategies for making multi-threaded systems where discussed in quite some detail.

#### Practical:

I saw a very cool JS program on emulating the functionality of `HashCode()` from Java on Stack Overflow and analyzing it was pretty fun and gave me a neat little function to use later on.

Code from Stack overflow:

```
String.prototype.hashCode = function(){
    var hash = 0;
    if (this.length == 0) return hash;
    for (i = 0; i < this.length; i++) {
        char = this.charCodeAt(i);
        hash = ((hash<<5)-hash)+char;
    }
}
```

```

        hash = hash & hash; // Convert to 32bit integer
    }
    return hash;
}

```

Firstly I learnt the pattern behind the whole thing .

$$s[0]*31^{(n-1)} + s[1]*31^{(n-2)} + \dots + s[n-1]$$

One cool thing in this code was that they had already bit shifted the code to avoid a multiply.

hash = ((hash<<5)-hash)+char;

instead of hash += char\*Math.pow(31,(lengthOfString -i));

That was really neat another improvement, not included in this original code , and which should help is to put the value of string length in an actual variable. This may seem like a childish way to code but I believe by storing that value, you would attain two things.

- a) Remove the need to compute the length again.
- b) Make it predictable for the processor to serialize the whole thing more efficiently.

I also found this column store based JS DB called Datavore on github

<https://github.com/StanfordHCL/datavore>

This is supposed to provide fast(under 100ms) access to over million data points while running in browser. This seemed quite underwhelming at first but then I tried to run it in browser to see how fast it was, and it was really good.

The project seems mostly abandoned with only a few commits, but it seems a great little project. The fact that they managed to do it without any other dependencies is pretty cool.

I still don't have any use for a column store DB but I learnt yet another tool in case I need it and I am going through the code of the whole thing and see if I can pick any cool tricks they used elsewhere.