Gaurav Singh | UFID: 36398850

COP 6726: Database Systems Implementation
Spring 2018
Weekly Assignment 8

20-03-2018:

- For column store the performance is 1000x faster than row store
- No SQL 2005-2015
- Not Only SQL hides behind scalability
- Many Column stores can process simple aggregates easily
- They fail in complicated datamining requirements
- Grokit was designed to specifically survive these situations
- Row Store DBs – no thought was given for user defined aggregates
- Row Store even if they support User Defined aggregates typically restrict it to one-page size whereas in Grokit they could 200 GB
- You are also limited to typically one return row
- If everything fails, do string processing in Hadoop
- Grokit takes 3 seconds to start a query
- Generalized Filters
    - Filter can be applied to a row
    - There is only a limited amount of fancy things you can do to a filter.
- Generalized Transformers
    - I want to take 4 column values and combine them in complicated ways and put the result in 5^{th} row.
    - It's hard in DS to stream this data at fast speed outside.
- GLADE Paper in 2009
- Generalized Input (GI)
    - Load in parallel
- GIST Generalized Iterative State Transformer
- Vectorized Processing

22-03-2018

- Iterator Model
- Lots of DBs use it, Its not fast, Zero Parallelization
- Operators can talk to each other with knowing each other at all.
- Group By is a blocking op, means you cant do other thing until the very end of group by
- Our project is a threadsafe way to communicate between two operators
- We are doing bottom up
- Its impossible to tell the spread if process in parallelism.
- MPP – Multiple processor parallelism
- MPP's biggest issue is you can virtually never do perfect balancing.